

# Exercise\_1

November 15, 2020

## 1 Topic 1: Exercise 1

### Importing the data

```
[25]: library(dplyr)
library(stringr)
library(PerformanceAnalytics)
library(foreach)
library(ggplot2)
library(reshape2)
library(MASS)
library(andrews)
```

```
[26]: d <- read.csv('.../.../datasets/Colleges.csv')
```

```
[27]: head(d)
```

	X <chr>	Private <chr>	Apps <dbl>	Accept <dbl>	Enroll <dbl>	Top10perc <dbl>	Top25perc <dbl>
A data.frame: 6 × 19	1 Abilene Christian University	Yes	1660	1232	721	23	52
	2 Adelphi University	Yes	2186	1924	512	16	29
	3 Adrian College	Yes	1428	1097	336	22	50
	4 Agnes Scott College	Yes	417	349	137	60	89
	5 Alaska Pacific University	Yes	193	146	55	16	44
	6 Albertson College	Yes	587	479	158	38	62

### 1.1 1- Identify the type of all variables

```
[28]: foreach (col = d, nms = colnames(d)) %do% {
  print(str_interp("${nms} is of type ${typeof(col)}"))
}
```

```
[1] "X is of type character"
[1] "Private is of type character"
[1] "Apps is of type double"
[1] "Accept is of type double"
[1] "Enroll is of type double"
[1] "Top10perc is of type double"
[1] "Top25perc is of type double"
```

```
[1] "F.Undergrad is of type double"
[1] "P.Undergrad is of type double"
[1] "Outstate is of type double"
[1] "Room.Board is of type double"
[1] "Books is of type double"
[1] "Personal is of type double"
[1] "PhD is of type double"
[1] "Terminal is of type double"
[1] "S.F.Ratio is of type double"
[1] "perc.alumni is of type double"
[1] "Expend is of type double"
[1] "Grad.Rate is of type integer"
```

1. 'X is of type character'
2. 'Private is of type character'
3. 'Apps is of type double'
4. 'Accept is of type double'
5. 'Enroll is of type double'
6. 'Top10perc is of type double'
7. 'Top25perc is of type double'
8. 'F.Undergrad is of type double'
9. 'P.Undergrad is of type double'
10. 'Outstate is of type double'
11. 'Room.Board is of type double'
12. 'Books is of type double'
13. 'Personal is of type double'
14. 'PhD is of type double'
15. 'Terminal is of type double'
16. 'S.F.Ratio is of type double'
17. 'perc.alumni is of type double'
18. 'Expend is of type double'
19. 'Grad.Rate is of type integer'

#### **represented as R type - numerical/categorical**

- X is of type character - categorical (id)
- Private is of type character - categorical binary
- Apps is of type integer - numerical discrete
- Accept is of type integer - numerical discrete

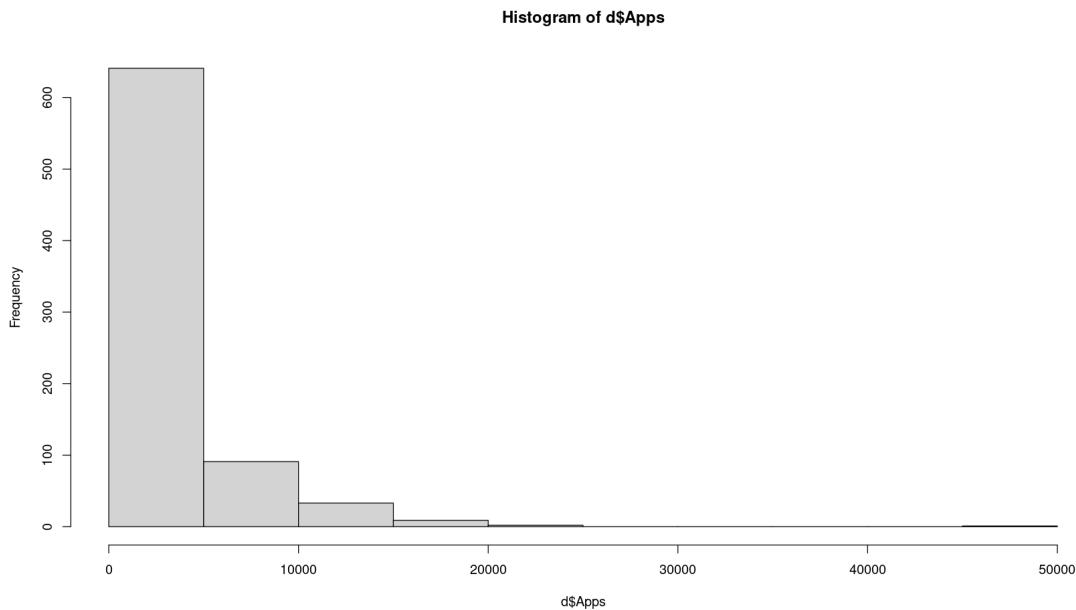
- Enroll is of type integer - numerical discrete
- Top10perc is of type integer - numerical discrete
- Top25perc is of type integer - numerical discrete
- F.Undergrad is of type integer - numerical discrete
- P.Undergrad is of type integer - numerical discrete
- Outstate is of type integer - numerical discrete
- Room.Board is of type integer - numerical discrete
- Books is of type integer - numerical discrete
- Personal is of type integer - numerical discrete
- PhD is of type integer - numerical discrete
- Terminal is of type integer - numerical discrete
- S.F.Ratio is of type double - numerical continuous
- perc.alumni is of type integer - numerical discrete
- Expend is of type integer - numerical discrete
- Grad.Rate is of type integer - numerical discrete

**1.2 2- Perform a visual analysis of each quantitative variable. Then do so considering the variable *private*. Describe each plot.**

### 1.2.1 Histograms

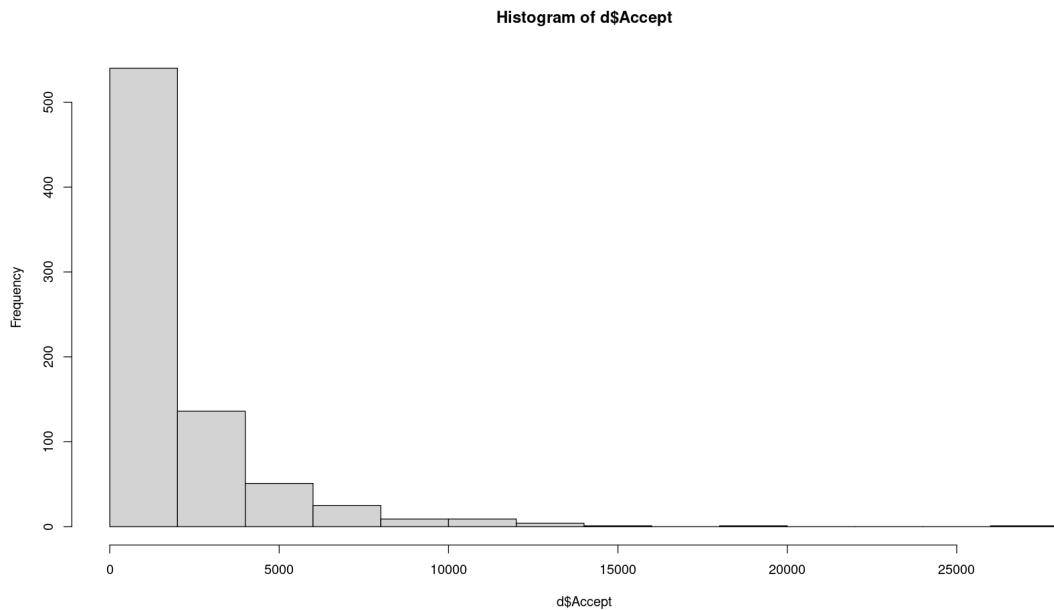
```
[29]: # setting plot sizes
options(repr.plot.width = 14, repr.plot.height = 8)
```

```
[30]: # Apps
hist(d$Apps)
```



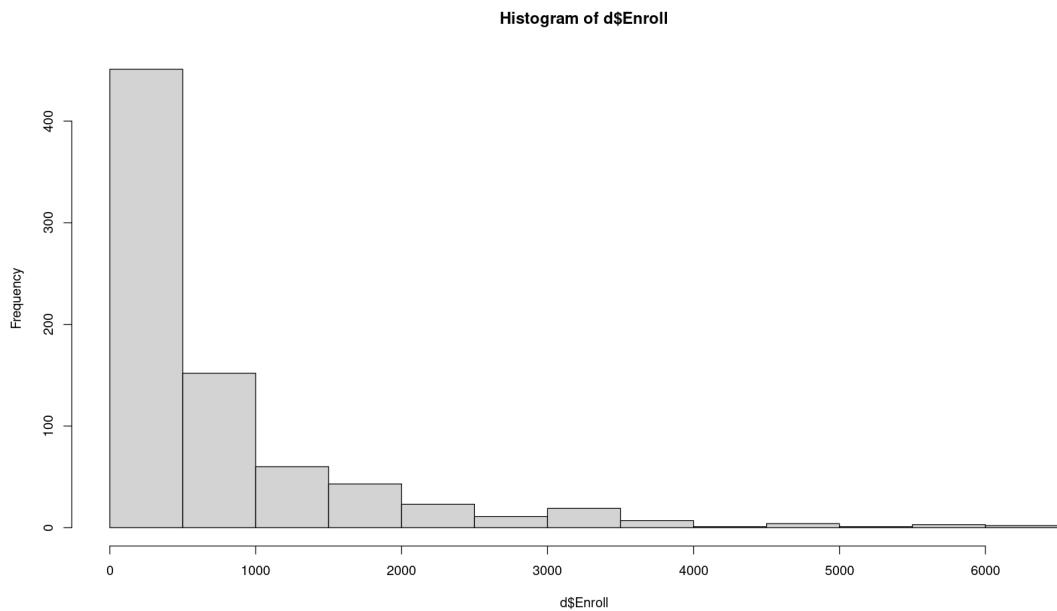
- Large right tail, higher frequency
- Long left tail
- Very right skewed
- Resembles gamma distribution

```
[31]: # Accept
hist(d$Accept)
```



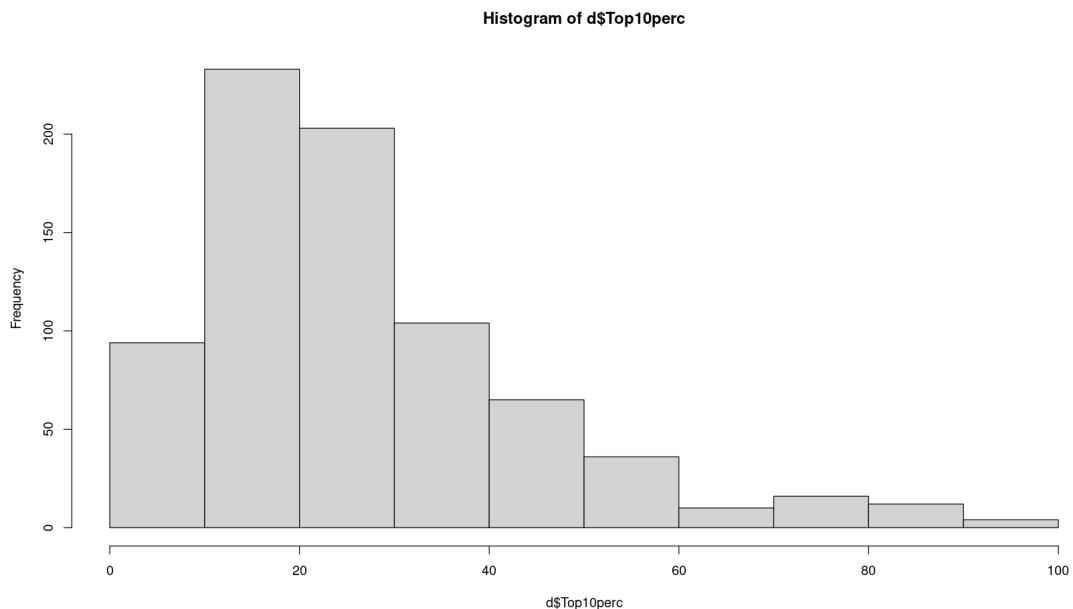
Pretty much the same as Apps but with a significantly shorter left tail. Probably because those students being accepted are just a subset of those that apply.

```
[32]: # Enroll
hist(d$Enroll)
```



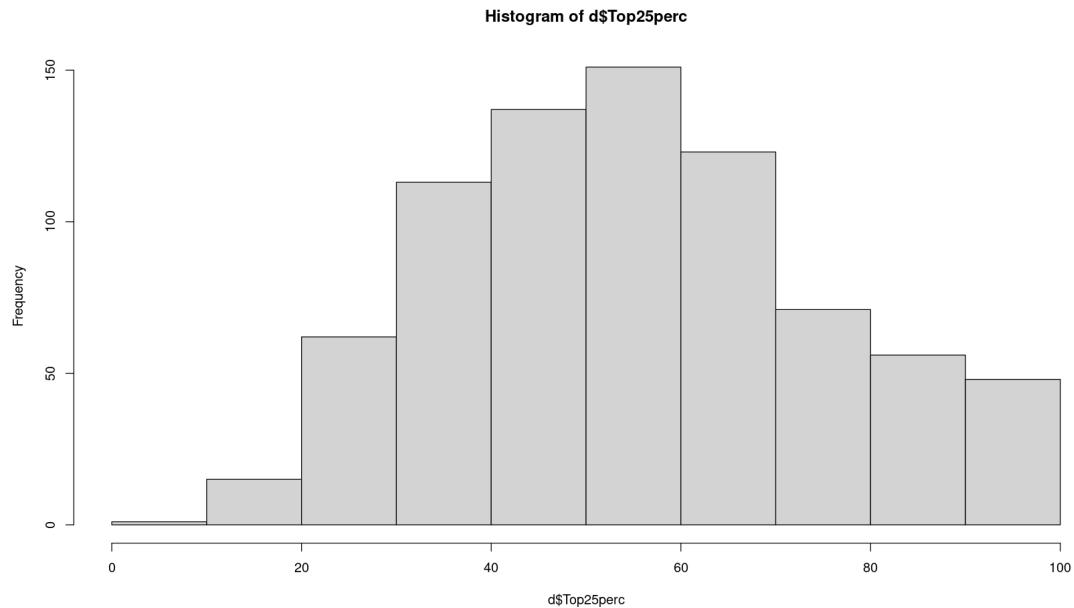
Enroll also tells the same story, but in this case our left tail is even smaller than the accept left tail. Maybe because significantly less students enroll than those that just get accepted

```
[33]: # Top10perc
hist(d$Top10perc)
```



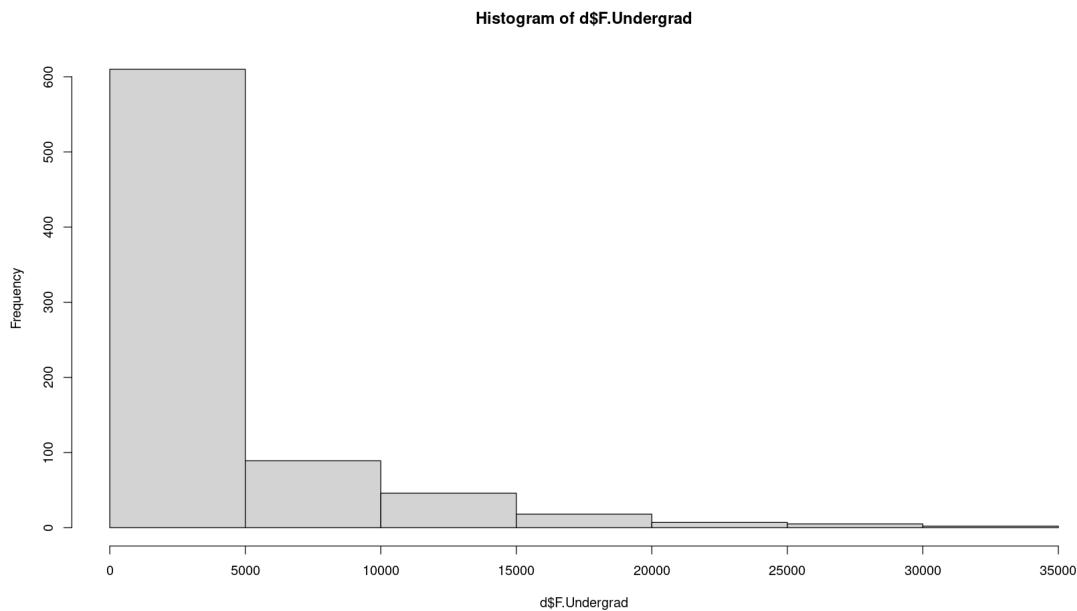
- long left tail
- short but high frequency right tail
- Resembles gaussian distribution
- most elements concentrate around 20

```
[34]: # Top25perc
hist(d$Top25perc)
```



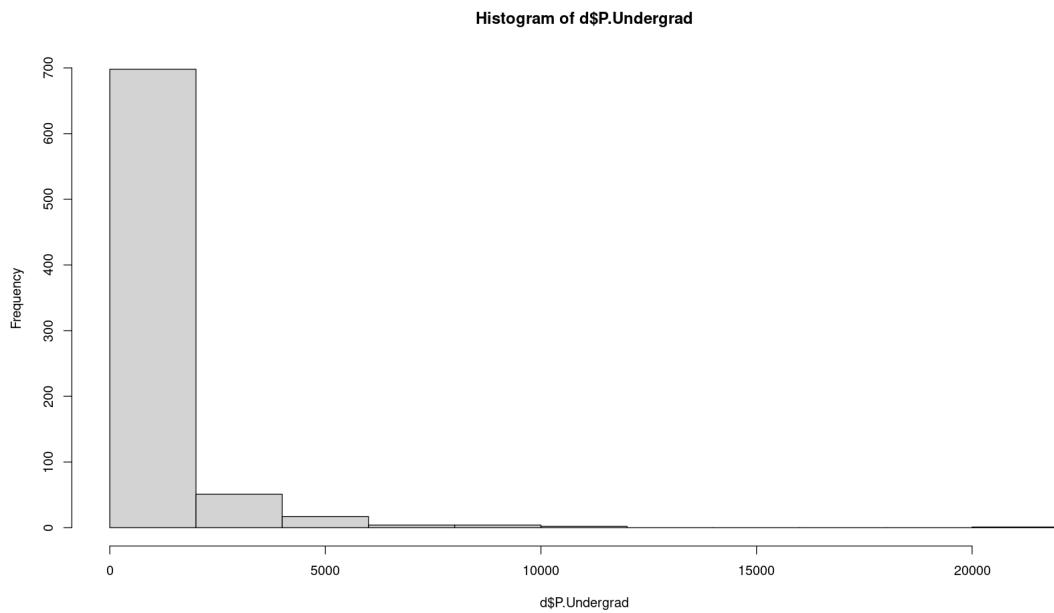
- mean is probably between 40-60
- small right tail
- large left tail
- Resembles gaussian distribution

```
[35]: # F.Undergrad
hist(d$F.Undergrad)
```



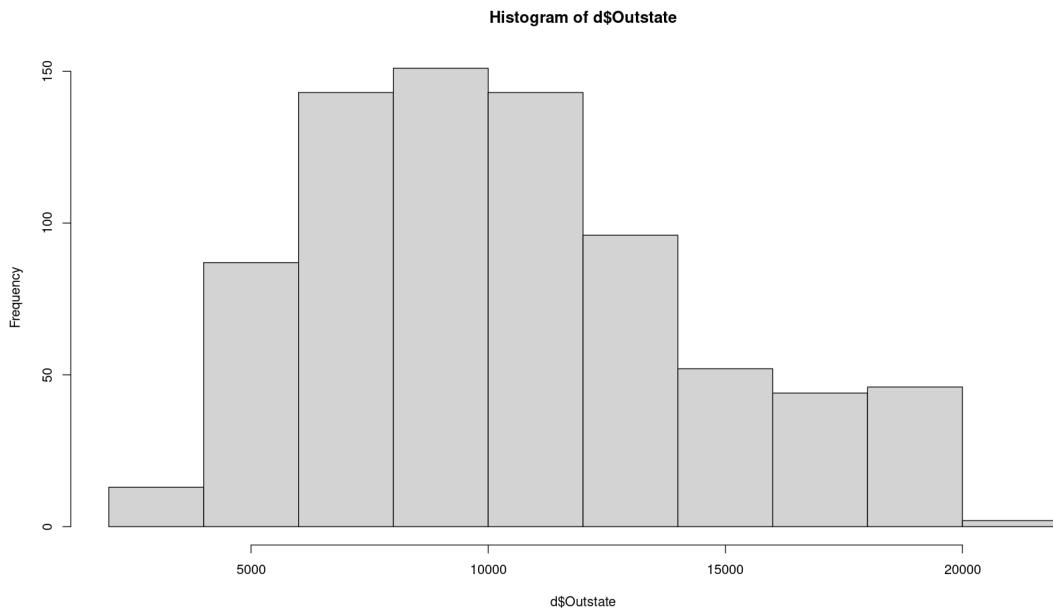
- long left tail
- most of the universities between 0-5000
- significantly less data at 5000+
- Resembles gamma distribution

```
[36]: # P.Undergrad
hist(d$P.Undergrad)
```



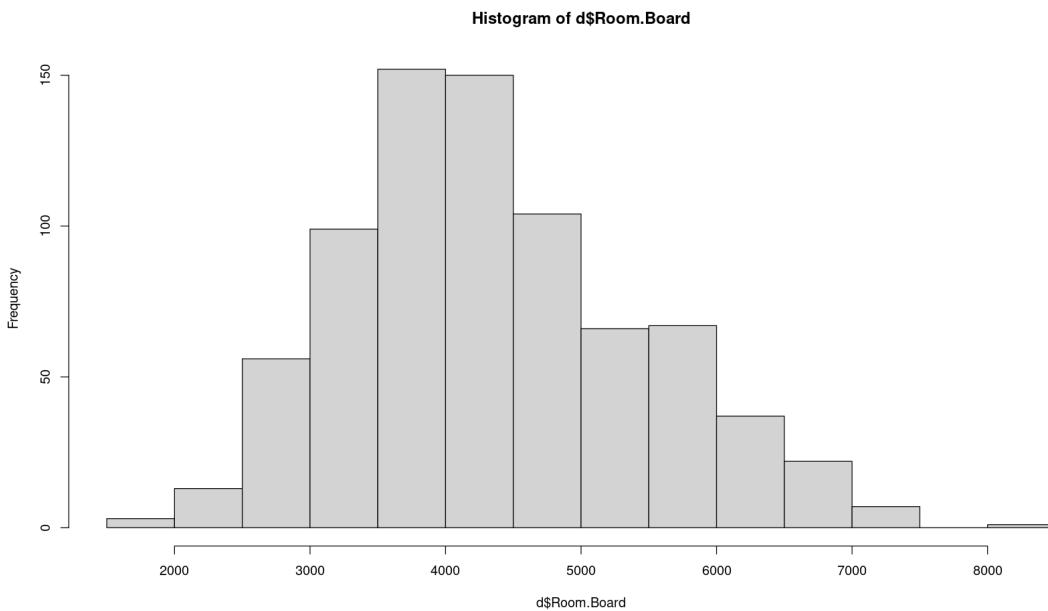
Similar to F.Undergrad we have a quite long right tail, this one resembles Apps more, and similar to F.Undergrad we also have a large frequency right tail and possibly more than 90% of the data between 0 and 5000

```
[37]: # Outstate  
hist(d$Outstate)
```



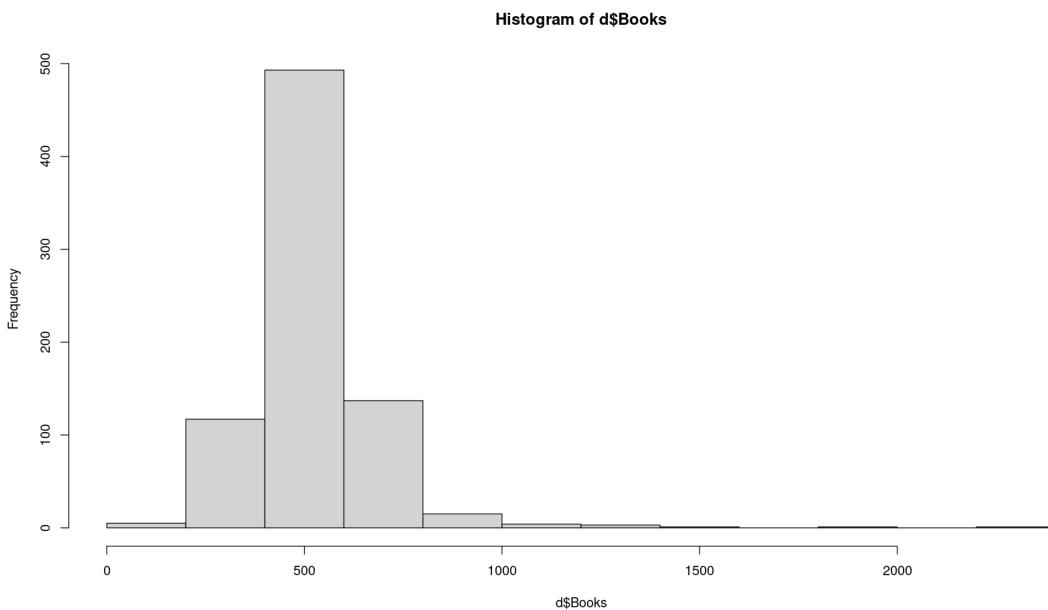
- Mean around 10000
- relatively heavy left tail but not very long
- light right tail
- resembles gaussian distribution

```
[38]: # Room.Board  
hist(d$Room.Board)
```



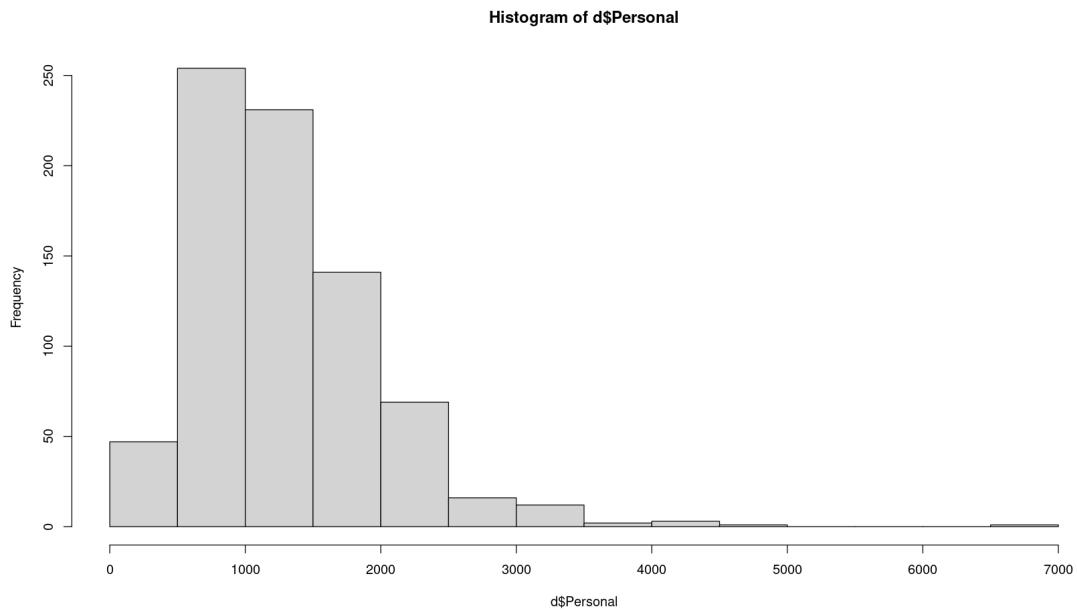
- resembles a slightly right skewed gaussian distribution
- relatively long left tail
- mean between 3500-4500

[39] : `# Books`  
`hist(d$Books)`



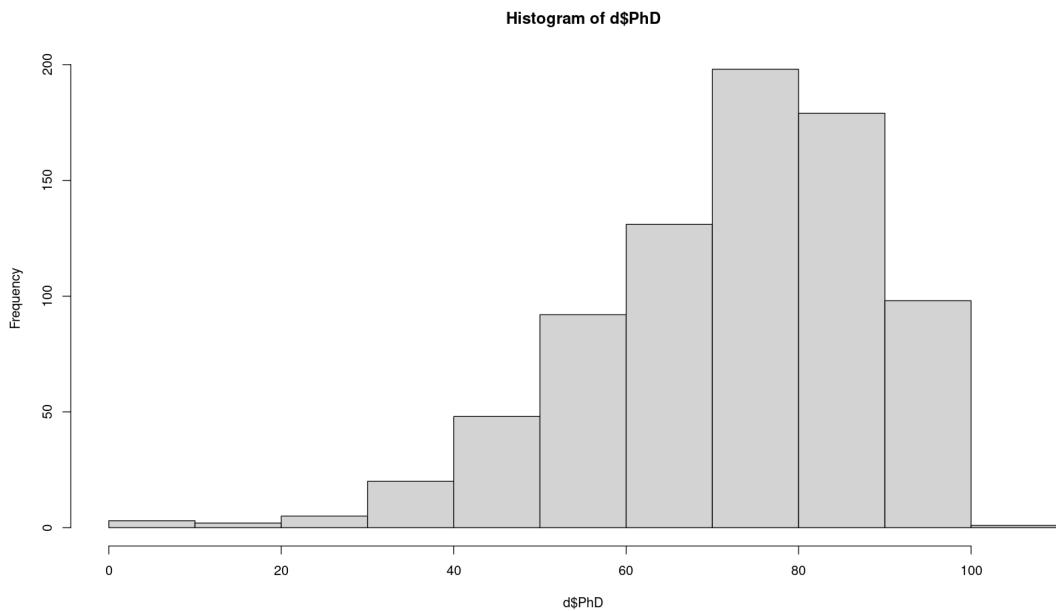
- long left tail
- very right skewed
- mean most probably around 500
- nearly all values between 400-750

```
[40]: # Personal  
hist(d$Personal)
```



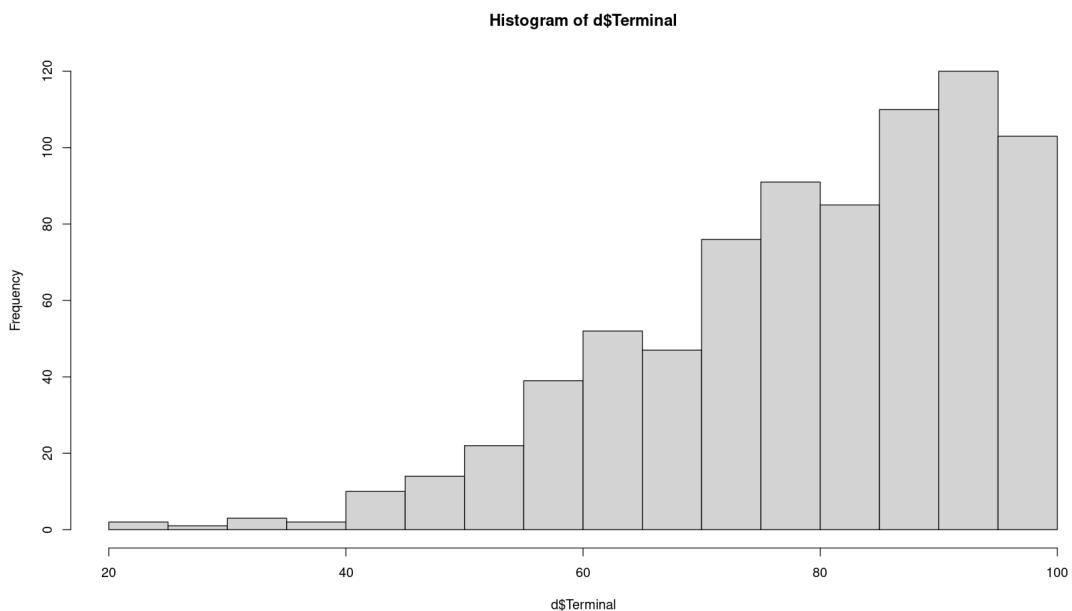
- very right skewed
- long left tail
- resembles a gamma distribution

```
[41]: # PhD  
hist(d$PhD)
```



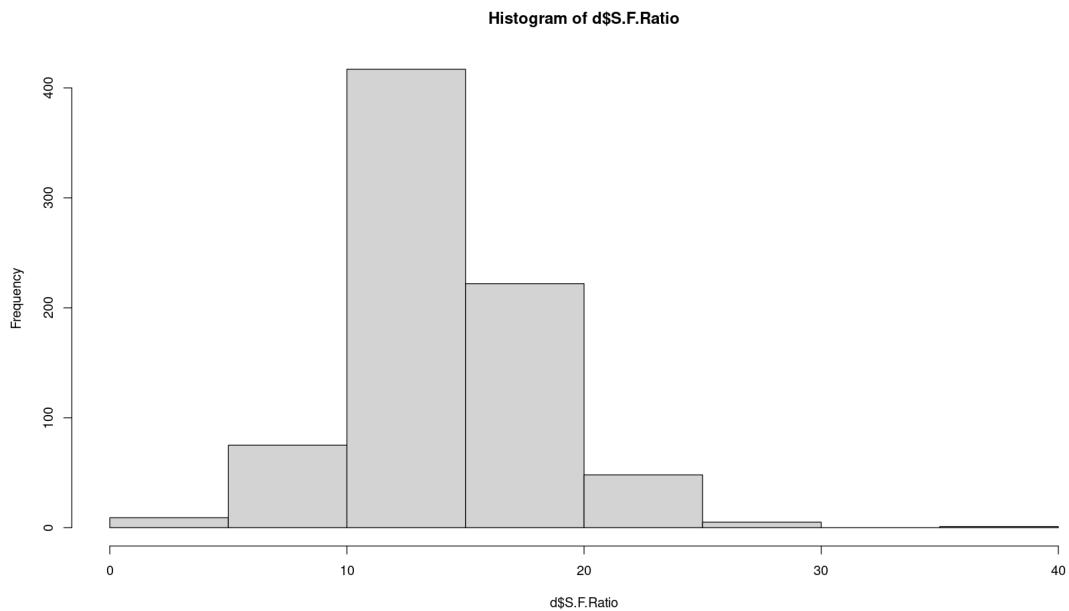
- long and light right tail
- very left skewed

```
[42]: # Terminal
hist(d$Terminal)
```



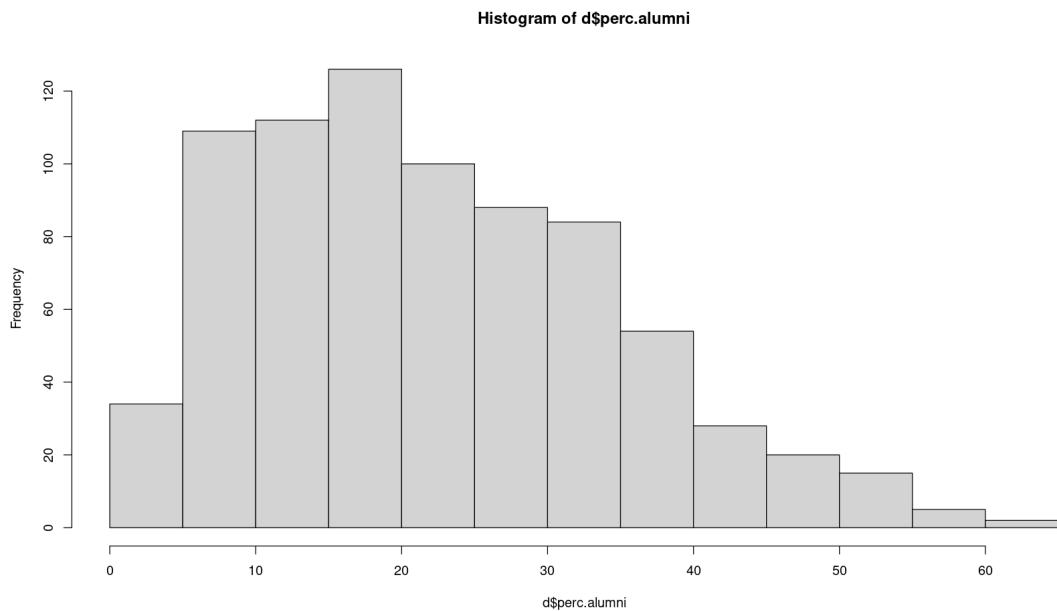
- very left skewed
- light and long right tail
- frequency increases with value (the higher the more common perhaps?)

```
[43]: # S.F.Ratio
hist(d$S.F.Ratio)
```



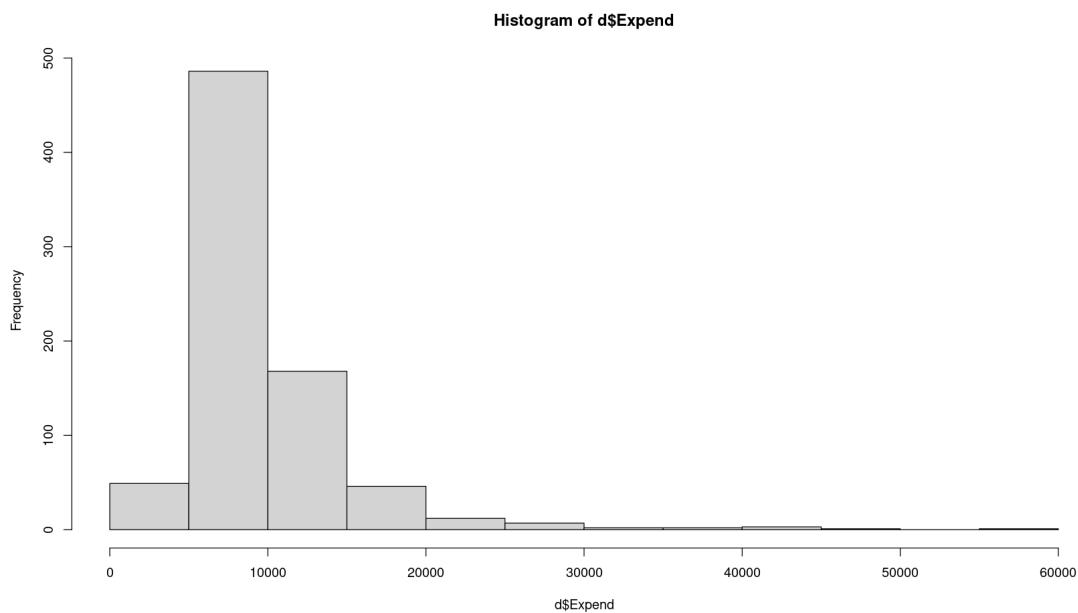
- most values between 10-20
- light and long left tail
- light right tail
- resembles gaussian distribution

```
[44]: # perc.alumni
hist(d$perc.alumni)
```



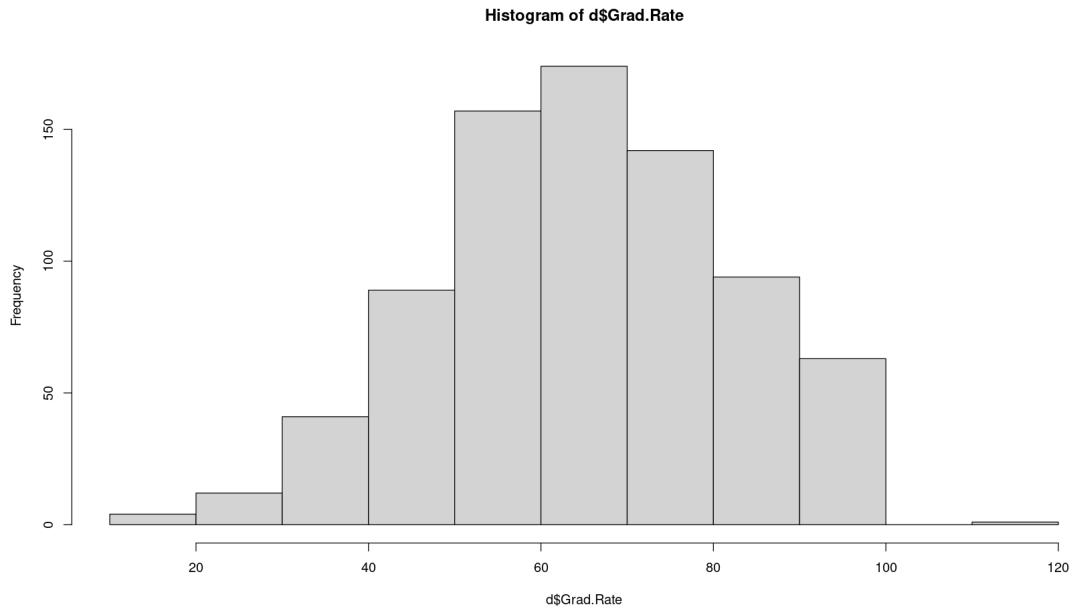
- right skewed
- long and light left tail
- relatively flat around the highest frequency part

```
[45] : # Expend
hist(d$Expend)
```



- very right skewed
- very long left tail
- most values around 10000
- resembles gamma distribution

```
[46]: # Grad.Rate
hist(d$Grad.Rate)
```



- quite resembles a gaussian distribution with mean between 50-70
- light leftmost bin frequency

### 1.2.2 Boxplots

```
[56]: # setting plot sizes
options(repr.plot.width = 18, repr.plot.height = 10)
```

```
[62]: # numeric columns
cols = colnames(d)[colnames(d) != 'X' & colnames(d) != 'Private']

# maximum per column
sapply(d %>% dplyr::select(cols), max, na.rm = TRUE)

# low val cols
low_val_c = c('Top10perc', 'Top25perc', 'PhD', 'Terminal', 'S.F.Ratio', 'perc.
  ↪alumni', 'Grad.Rate')
low_val = melt(d, id.vars='X', measure.vars=low_val_c)
```

```

# med val cols
med_val_c = c('Enroll', 'Room.Board', 'Books', 'Personal')
med_val = melt(d, id.vars='X', measure.vars=med_val_c)

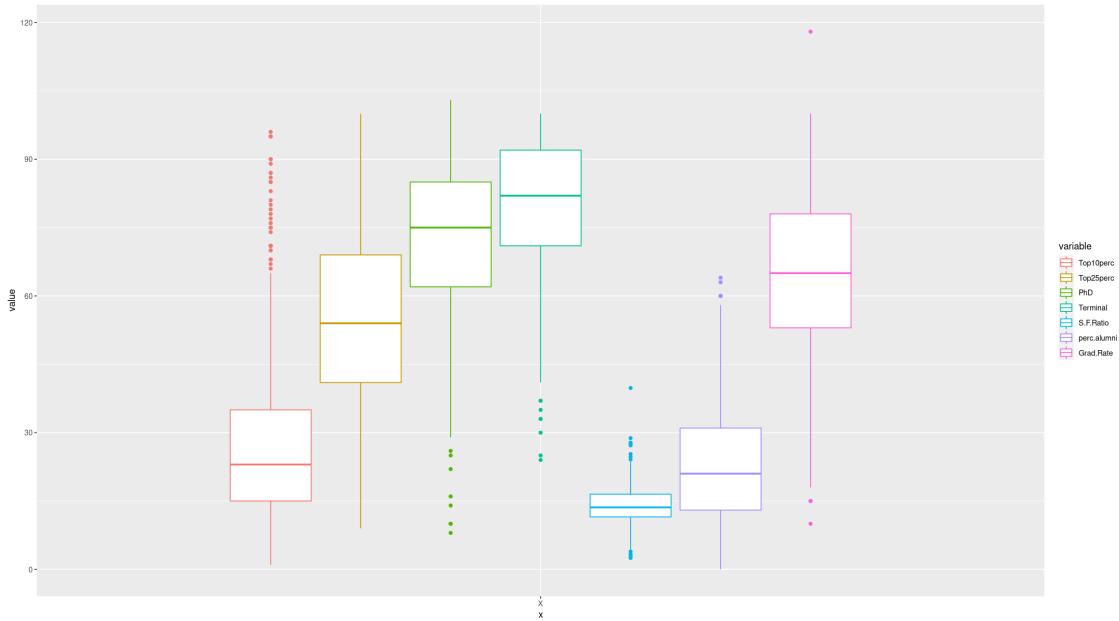
# high val cols
high_val_c = c('Apps', 'Accept', 'F.Undergrad', 'P.Undergrad', 'Outstate', 'Expend')
high_val = melt(d, id.vars='X', measure.vars=high_val_c)

```

Apps 48094 Accept 26330 Enroll 6392 Top10perc 96 Top25perc 100 F.Undergrad 31643  
 P.Undergrad 21836 Outstate 21700 Room.Board 8124 Books 2340 Personal 6800 PhD 103  
 Terminal 100 S.F.Ratio 39.8 perc.alumni 64 Expend 56233 Grad.Rate 118

Variables with low values (maximum under 120)

```
[63]: low_val %>%
  ggplot(aes(x="X", y=value)) +
  geom_boxplot(aes(color=variable))
```



```
[108]: # quantiles
sapply(d %>% dplyr::select(low_val_c), quantile, na.rm = TRUE)
```

Note: Using an external vector in selections is ambiguous.

Use `all\_of(low\_val\_c)` instead of `low\_val\_c` to silence this message.

See <<https://tidyselect.r-lib.org/reference/faq-external-vector.html>>.

This message is displayed once per session.

		Top10perc	Top25perc	PhD	Terminal	S.F.Ratio	perc.alumni	Grad.I
	0%	1	9	8	24	2.5	0	10
A matrix: 5 × 7 of type dbl	25%	15	41	62	71	11.5	13	53
	50%	23	54	75	82	13.6	21	65
	75%	35	69	85	92	16.5	31	78
	100%	96	100	103	100	39.8	64	118

For the low value variables we can notice a few things:

- Top10perc

In our boxplot we can see that the median for this variable is 23, the maximum is 96 and the minimum is 1 with the 25% and 75% percentiles sitting at 15 and 35 respectively.

The range for this variable is of 95 and the IQR (interquartile range) is equal to 20.

- Top25perc

In our boxplot we can see that the median for this variable is 54, the maximum is 100 and the minimum is 9 with the 25% and 75% percentiles sitting at 41 and 69 respectively.

The range for this variable is of 91 and the IQR is equal to 28

- PhD

In our boxplot we can see that the median for this variable is 75, the maximum is 103 and the minimum is 8 with the 25% and 75% percentiles sitting at 62 and 85 respectively.

The range for this variable is of 95 and the IQR is equal to 23

- Terminal

In our boxplot we can see that the median for this variable is 82, the maximum is 100 and the minimum is 24 with the 25% and 75% percentiles sitting at 71 and 92 respectively.

The range for this variable is of 76 and the IQR is equal to 21

- S.F.Ratio

In our boxplot we can see that the median for this variable is 13.6, the maximum is 39.8 and the minimum is 2.5 with the 25% and 75% percentiles sitting at 11.5 and 16.5 respectively.

The range for this variable is of 37.3 and the IQR is equal to 5

- perc.alumni

In our boxplot we can see that the median for this variable is 21, the maximum is 64 and the minimum is 0 with the 25% and 75% percentiles sitting at 13 and 31 respectively.

The range for this variable is of 64 and the IQR is equal to 18

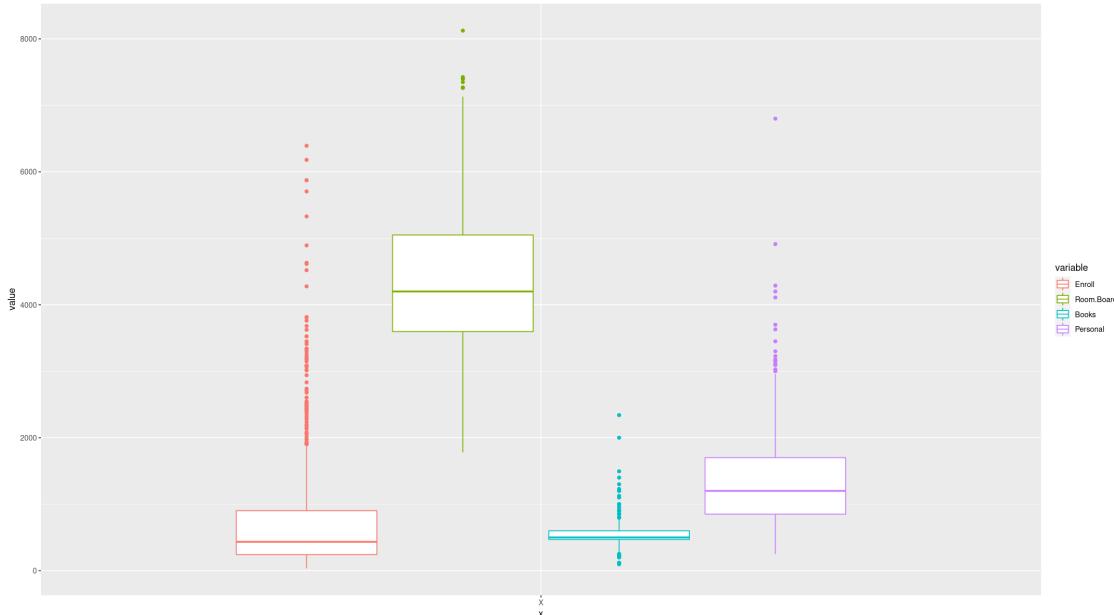
- Grad.Rate

In our boxplot we can see that the median for this variable is 65, the maximum is 118 and the minimum is 10 with the 25% and 75% percentiles sitting at 53 and 78 respectively.

The range for this variable is of 108 and the IQR is equal to 25

Variables with medium values (maximum less than or equal to 8124 and not in low values list)

```
[65]: med_val %>%
  ggplot(aes(x="X", y=value)) +
  geom_boxplot(aes(color=variable))
```



```
[69]: # quantiles
sapply(d %>% dplyr::select(med_val_c), quantile, na.rm = TRUE)
```

	Enroll	Room.Board	Books	Personal
0%	35	1780	96	250
25%	242	3597	470	850
50%	434	4200	500	1200
75%	902	5050	600	1700
100%	6392	8124	2340	6800

For the medium value variables we can notice a few things:

- Enroll

In our boxplot we can see that the median for this variable is 494, the maximum is 6392 and the minimum is 35 with the 25% and 75% percentiles sitting at 242 and 902 respectively.

The range for this variable is of 6357 and the IQR (interquartile range) is equal to 660.

- Room.Board

In our boxplot we can see that the median for this variable is 4200, the maximum is 8124 and the minimum is 1780 with the 25% and 75% percentiles sitting at 3597 and 5050 respectively.

The range for this variable is of 6344 and the IQR is equal to 1453.

- Books

In our boxplot we can see that the median for this variable is 500, the maximum is 2340 and the minimum is 96 with the 25% and 75% percentiles sitting at 470 and 600 respectively.

The range for this variable is of 2244 and the IQR is equal to 130

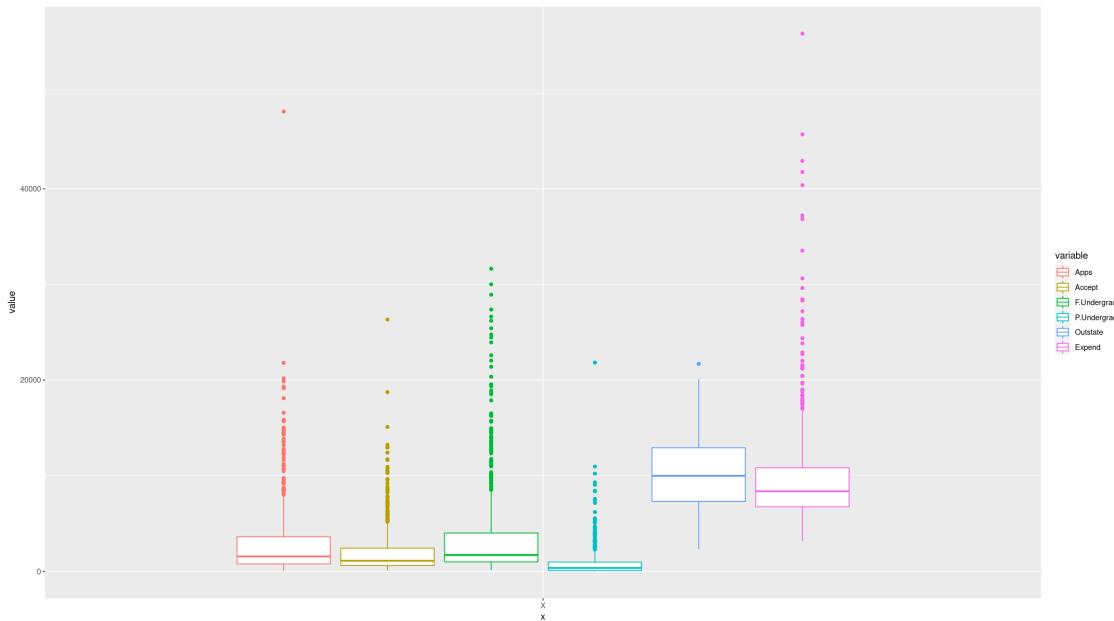
- Personal

In our boxplot we can see that the median for this variable is 1200, the maximum is 6800 and the minimum is 250 with the 25% and 75% percentiles sitting at 850 and 1700 respectively.

The range for this variable is of 6550 and the IQR is equal to 850

### Variables with high values (maximums above 8124)

```
[70]: high_val %>%
  ggplot(aes(x="X", y=value)) +
  geom_boxplot(aes(color=variable))
```



```
[72]: # quantiles
sapply(d %>% dplyr::select(high_val_c), quantile, na.rm = TRUE)
```

	Apps	Accept	F.Undergrad	P.Undergrad	Outstate	Expend
0%	81	72	139	1	2340	3186
25%	776	604	992	95	7320	6751
50%	1558	1110	1707	353	9990	8377
75%	3624	2424	4005	967	12925	10830
100%	48094	26330	31643	21836	21700	56233

For the high value variables we can notice a few things:

- Apps

In our boxplot we can see that the median for this variable is 1558, the maximum is 48094 and the minimum is 81 with the 25% and 75% percentiles sitting at 776 and 3624 respectively.

The range for this variable is of 48013 and the IQR (interquartile range) is equal to 2848.

- Accept

In our boxplot we can see that the median for this variable is 1110, the maximum is 26330 and the minimum is 72 with the 25% and 75% percentiles sitting at 604 and 2424 respectively.

The range for this variable is of 26258 and the IQR is equal to 1820.

- F.Undergrad

In our boxplot we can see that the median for this variable is 1707, the maximum is 31643 and the minimum is 139 with the 25% and 75% percentiles sitting at 992 and 4005 respectively.

The range for this variable is of 31504 and the IQR is equal to 3013

- P.Undergrad

In our boxplot we can see that the median for this variable is 353, the maximum is 21836 and the minimum is 1 with the 25% and 75% percentiles sitting at 95 and 967 respectively.

The range for this variable is of 21835 and the IQR is equal to 872

- Outstate

In our boxplot we can see that the median for this variable is 9990, the maximum is 21700 and the minimum is 2340 with the 25% and 75% percentiles sitting at 7320 and 12925 respectively.

The range for this variable is of 19360 and the IQR is equal to 5605

- Expend

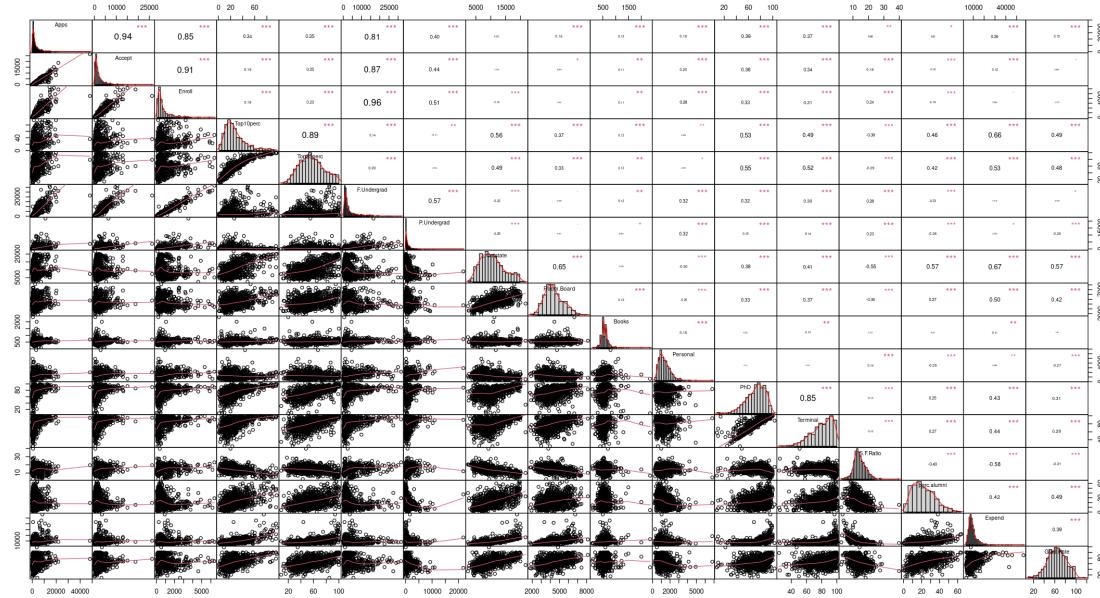
In our boxplot we can see that the median for this variable is 8377, the maximum is 56233 and the minimum is 3186 with the 25% and 75% percentiles sitting at 6751 and 10830 respectively.

The range for this variable is of 53047 and the IQR is equal to 4079

### 1.3 3- Perform a visual analysis of all quantitative variables together. Then, Perform a visual analysis of each of the quantitative variables taking into account the variable Private.

#### 1.3.1 Correlation and scatterplots for all variables

```
[74]: pa <- d %>% dplyr::select(cols)
chart.Correlation(pa, histogram=TRUE, pch=19, method="pearson")
```



Here we can see a plot with all the numerical variables, a histogram and density plot in the main diagonal. To the bottom of the diagonal we see scatterplots for each pair of variables and a line that goes through the scatterplot. Above the main diagonal we see all the correlation coefficients for each pair of variables.

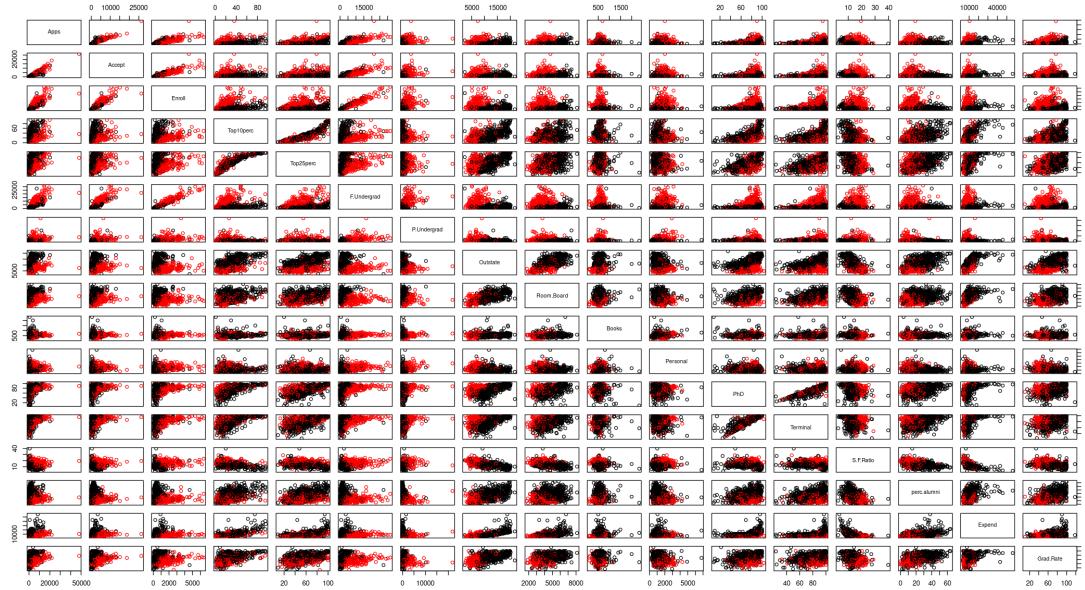
We see there are some highly correlated variables and some others which have negligible correlation. All correlation coefficients are Pearson correlation coefficients.

The variables with the highest correlation are the following:

- F.Undergrad vs Enroll with very high correlation of 0.96
- Apps vs Accept with a similarly high correlation coefficient of 0.94
- Accept vs Enroll also have a very high correlation of 0.91
- Top10perc vs Top25perc have a correlation of 0.89
- F.Undergrad vs Accept have a correlation of 0.
- Both Apps vs Enroll and Phd vs Terminal have a correlation of 0.85

```
[100]: # setting colors
color_1 <- "black"
color_2 <- "red"
colors <- as.numeric(as.factor(d$Private)) - 1
i <- 1
for (num in colors) {
  colors[i] = ifelse(num==1, color_1, color_2)
  i <- i + 1
}
```

```
[101]: pa <- d %>% dplyr::select(cols)
pairs(pa,pch=1,col=colors)
```



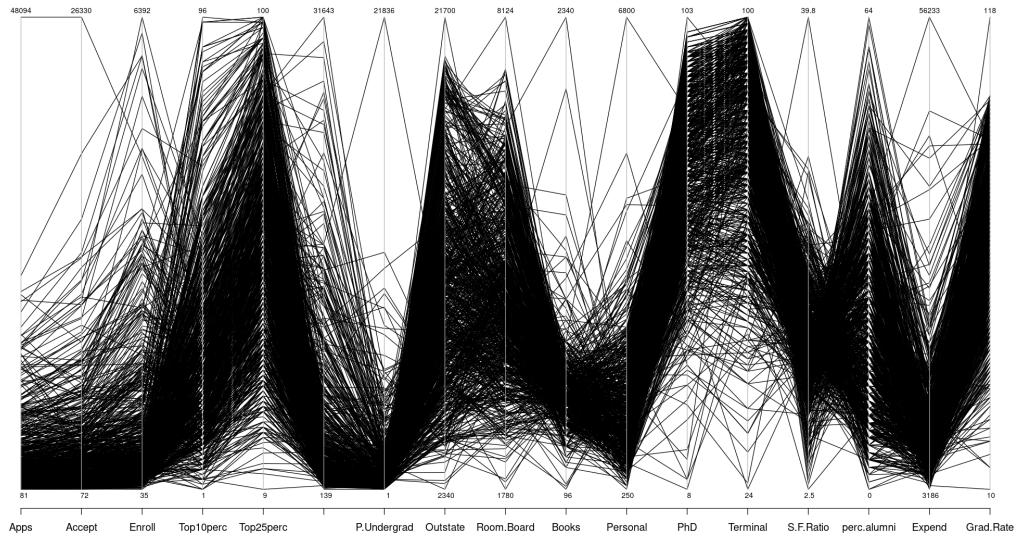
In most of these scatter plots we see an interesting pattern, where for variables plotted versus Top10perc, Top20perc, PhD and Terminal there seems to be no clear divide between each one ofg the groups, except when compared to variables that significantly diffentiate both groups.

The clear differentiator variables seem to be P.Undergrad, F.Undergrad, S.F.Ratio, Apps, Accept, Enrool and outstate.

And sure, I could definitely agree that the size of each plot might be overestimating such division in groups, but some of these variables do show a stark difference in values.

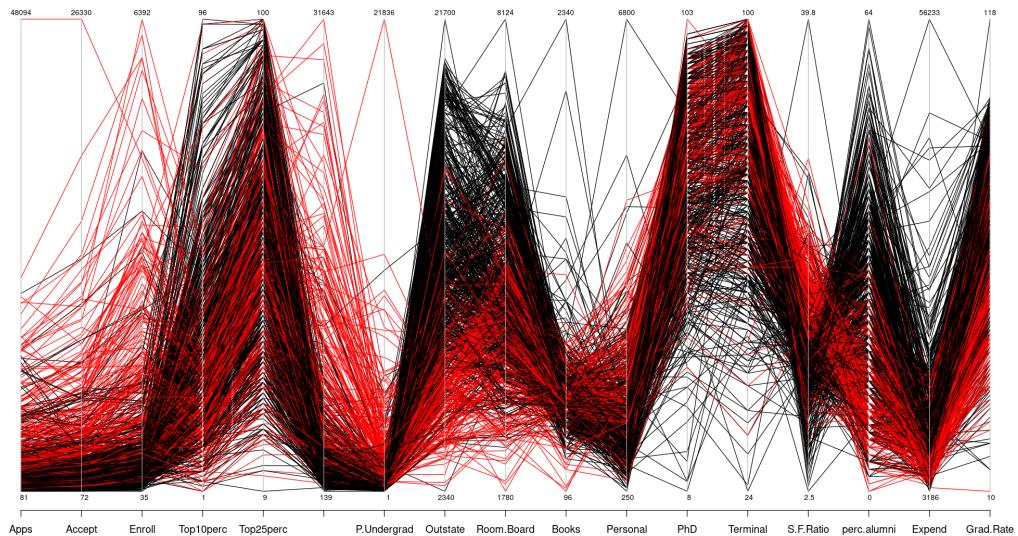
### 1.3.2 PCP

```
[102]: parcoord(pa,var.label=TRUE)
```



We can basically just see the spread of the data here, the only thing I take from it is that apps and accept have really very few outliers in terms of quantity, but they're significantly larger than the rest. Shows how skewed these variables are.

```
[103]: parcoord(pa,var.label=TRUE, col=colors)
```



We can see that the more extreme variables of Apps, Accept, P.Undergrad, F.Undergrad seem to

be affected by whether it's private or not.

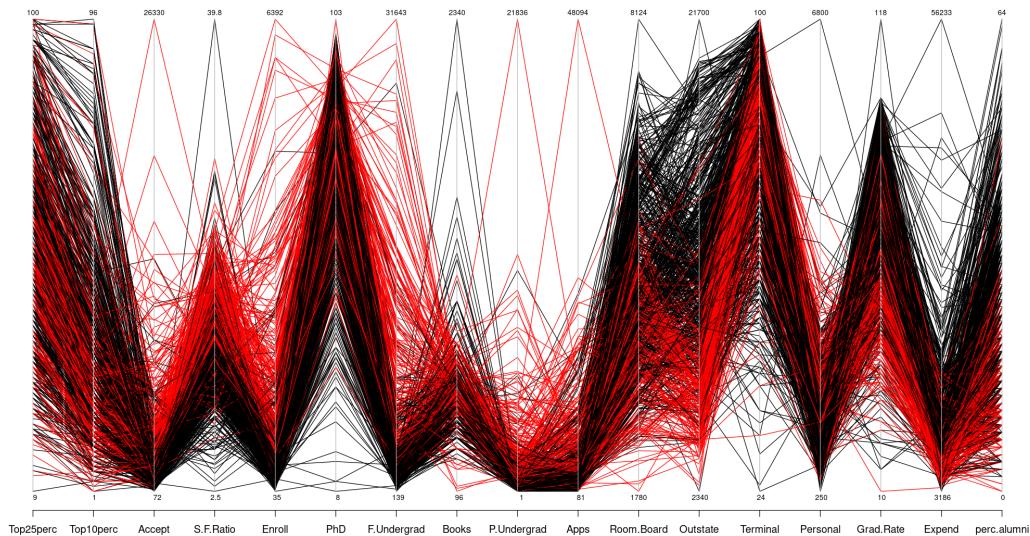
In the case of outstate too, but instead of being the highest values, the values where it's private correspond to the lowest values.

For Grad.Rate it also seem somewhat divided and in a similar way for perc.alumni.

For Top10perc, Top25perc, PhD or Terminal there seems to be no noticeable difference made by whether it's private or not.

**Randomizing the order of the variables we get a different but perhaps better PCP.**

```
[107]: cols <- sample(cols)
pa <- pa %>% dplyr::select(cols)
parcoord(pa, var.label=TRUE, col=colors)
```



Here we can more clearly see that for S.F.Ratio there is a stark difference between both groups, where the group labeled as Private seems to have the higher values.

For the rest of the variables we get much of the same we explained earlier.

**If we were to line up variables with similarly shaped histograms, we could get an interesting PCP, let's try that.** We will exclude Top10perc and Top25perc as these do not contribute too much to the plot

```
[121]: # right skewed variables
right_skewed <- c('Apps', 'Accept', 'P.Undergrad', 'F.Undergrad', 'Personal', 'Expend', 'Books', 'perc.alumni')
right_skewed <- d %>% dplyr::select(right_skewed)
```

```

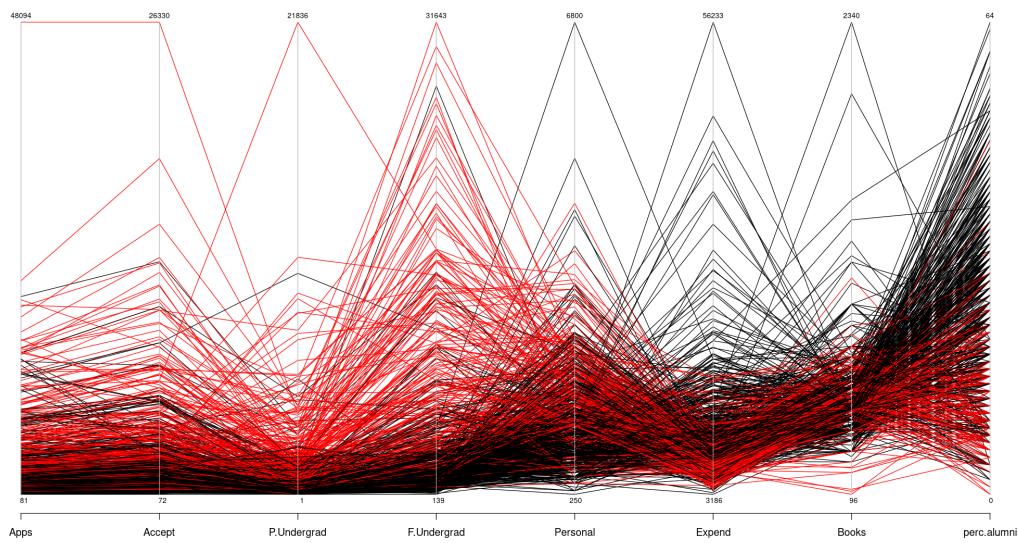
# normal-like variables
normal_like <- c('Outstate', 'Room.Board', 'perc.alumni', 'Grad.Rate', 'S.F.Ratio')
normal_like <- d %>% dplyr::select(normal_like)

# left skewed variables
left_skewed <- c('PhD', 'Terminal')
left_skewed <- d %>% dplyr::select(left_skewed)

```

### right skewed variables PCP

```
[122]: parcoord(right_skewed, var.label=TRUE, col=colors)
```

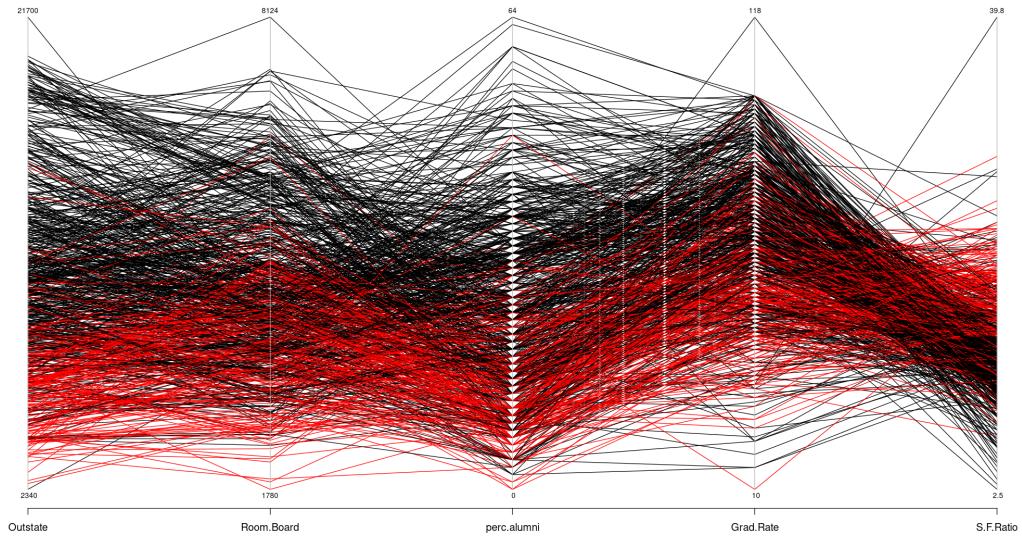


After grouping similarly shaped variables the plots are somewhat clearer. Confirming most of our previous observations.

The general idea is that for these variables, pretty much everything but perc.alumni, Books and Expend follow a pattern where the higher-than-usual variables pretty much belong completely to entries of Private colleges.

### normal-like variables PCP

```
[120]: parcoord(normal_like, var.label=TRUE, col=colors)
```

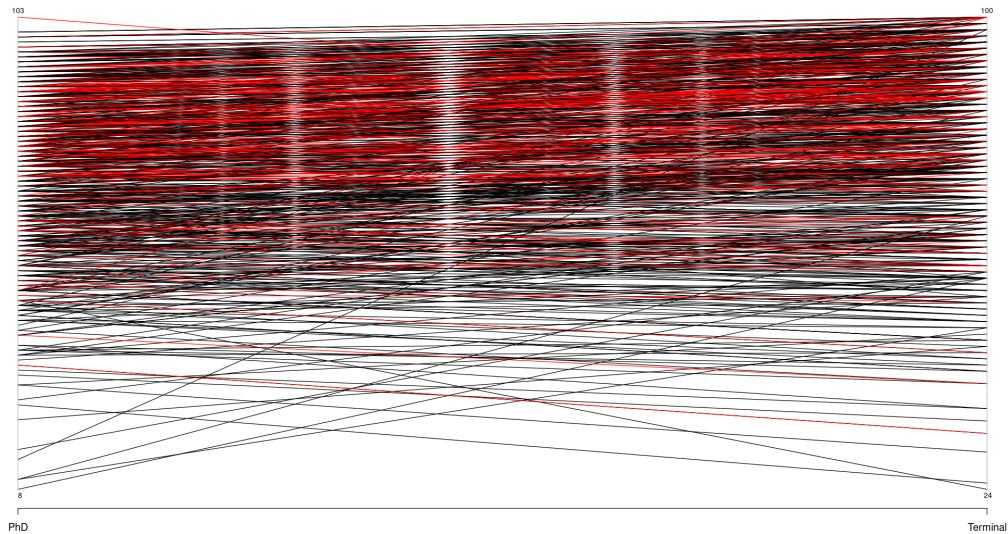


I particularly like this one. These variables look significantly clearer here than in the first PCP.

The first observations still hold though. For pretty much all these variables there's a significant divide in the same direction except for S.F.Ratio where that pattern is inverted.

#### left skewed variables PCP

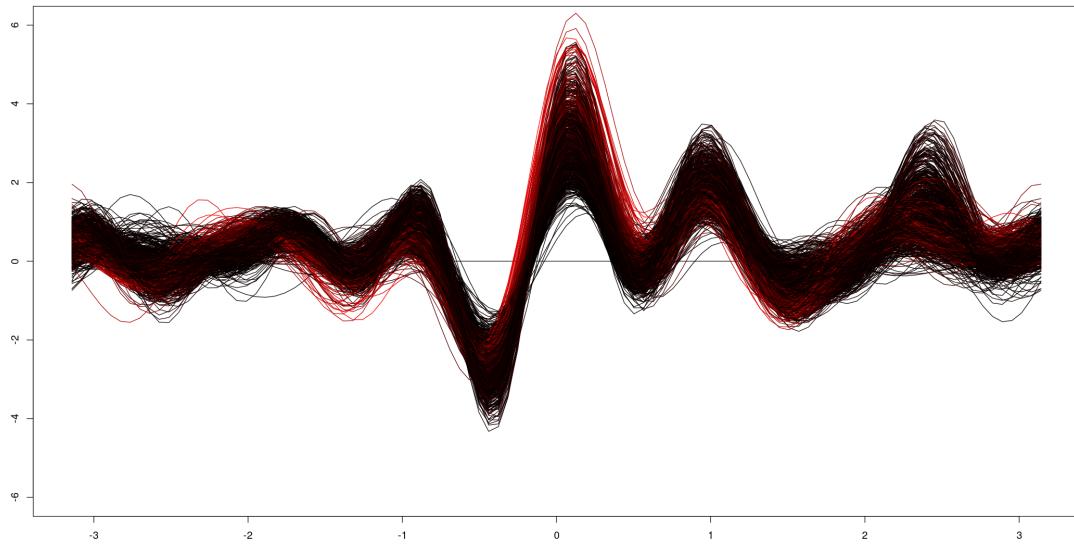
```
[117]: parcoord(left_skewed,var.label=TRUE, col=colors)
```



being private or not makes no difference here, therefore not very interesting.

### 1.3.3 Andrews' plot

```
[124]: par(mfrow=c(1,1))
andrews(as.data.frame(cbind(pa,as.factor(d$Private))),clr=7,ymax=6)
```



The Andrews' plot makes it difficult to appreciate a significant difference between both groups in my opinion. Perhaps I might need to read up more on it but it doesn't show tell me too much about the values other than around the middle. Where we have the 0 and around -2, where the groups seem more divided than everywhere else, taking different directions in the dataset.

```
[ ]:
```